

# Database Enhancement Narrative

---

## Travlr Getaways Polyglot Persistence and Security Implementation

---

### Artifact Description

The database enhancement represents a comprehensive transformation of the Travlr Getaways application's data layer, implementing polyglot persistence architecture with PostgreSQL and MongoDB. This enhancement was completed for CS 499 Milestone Four to demonstrate advanced database design principles, security implementation, and enterprise-grade data management practices.

The original application used only MongoDB for all data storage, which, while suitable for content management, lacked the transactional guarantees and security features necessary for handling sensitive user data and financial transactions. The enhancement introduces a sophisticated dual-database architecture that leverages the strengths of both relational and document databases.

### Justification for Inclusion

I selected this enhancement for my ePortfolio because it demonstrates my mastery of database design principles and their practical application in enterprise environments. The implementation of polyglot persistence showcases my understanding of different database paradigms and their appropriate use cases.

This enhancement is particularly valuable because it:

1. **Demonstrates database design expertise** through strategic data distribution
2. **Shows security implementation skills** with row-level security and audit logging
3. **Exhibits performance optimization** through indexing and query optimization
4. **Addresses real-world scalability challenges** with connection pooling and backup systems

The enhancement process showcases my ability to design, implement, and secure complex database systems that meet enterprise requirements for data integrity, security, and performance.

### Enhancement Process and Learning

#### Database Architecture Analysis

#### Original Architecture Assessment:

- **Single Database:** MongoDB only for all data types
- **Security Concerns:** No row-level security or data isolation
- **Transaction Limitations:** No ACID compliance for critical operations
- **Backup Strategy:** Basic MongoDB backup without point-in-time recovery
- **Performance Issues:** Missing indexes and unoptimized queries

#### Strategic Data Distribution Plan:

1. **PostgreSQL:** Transactional data requiring ACID compliance (bookings, payments, users)
2. **MongoDB:** Content data requiring flexibility (trips, descriptions, images)
3. **Data Synchronization:** Event-driven updates between databases
4. **Security Implementation:** Row-level security and audit logging

#### Polyglot Persistence Implementation

##### PostgreSQL Schema Design:

```
-- Users table with enhanced security
CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    username VARCHAR(50) UNIQUE NOT NULL,
    email VARCHAR(100) UNIQUE NOT NULL,
    password_hash VARCHAR(255) NOT NULL,
    salt VARCHAR(255) NOT NULL,
    role VARCHAR(20) DEFAULT 'user',
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    last_login TIMESTAMP,
    is_active BOOLEAN DEFAULT true
);

-- Bookings table with foreign key constraints
CREATE TABLE bookings (
    id SERIAL PRIMARY KEY,
    user_id INTEGER NOT NULL REFERENCES users(id) ON DELETE CASCADE,
    trip_code VARCHAR(20) NOT NULL,
    booking_date TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    travel_date DATE NOT NULL,
    travelers INTEGER NOT NULL CHECK (travelers > 0),
    total_amount DECIMAL(10,2) NOT NULL CHECK (total_amount > 0),
    status VARCHAR(20) DEFAULT 'pending',
    payment_reference VARCHAR(100),
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

-- Audit log table for tracking all changes
```

```

CREATE TABLE audit_log (
    id SERIAL PRIMARY KEY,
    table_name VARCHAR(50) NOT NULL,
    record_id INTEGER NOT NULL,
    operation VARCHAR(10) NOT NULL,
    old_values JSONB,
    new_values JSONB,
    user_id INTEGER,
    timestamp TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    ip_address INET,
    user_agent TEXT
);

```

### **MongoDB Schema Optimization:**

```

// Enhanced trip schema with indexing
const tripSchema = new mongoose.Schema({
  code: { type: String, required: true, unique: true, index: true },
  name: { type: String, required: true, text: true },
  length: { type: Number, required: true, min: 1 },
  start: { type: Date, required: true, index: true },
  resort: { type: String, required: true, text: true },
  perPerson: { type: Number, required: true, min: 0 },
  image: { type: String, required: true },
  description: { type: String, required: true, text: true },
  // Search optimization fields
  searchKeywords: [String],
  priceRange: { type: String, enum: ['budget', 'mid-range', 'luxury'] },
  // Metadata for performance
  createdAt: { type: Date, default: Date.now, index: true },
  updatedAt: { type: Date, default: Date.now },
  isActive: { type: Boolean, default: true, index: true }
}, {
  // Text search index
  text: {
    name: 'text',
    resort: 'text',
    description: 'text',
    searchKeywords: 'text'
  }
});

```

### **Row-Level Security Implementation**

#### **PostgreSQL Security Policies:**

```

-- Enable row-level security
ALTER TABLE bookings ENABLE ROW LEVEL SECURITY;
ALTER TABLE users ENABLE ROW LEVEL SECURITY;

-- User data isolation policy
CREATE POLICY user_data_isolation ON bookings
  FOR ALL TO authenticated_user
  USING (user_id = current_setting('app.current_user_id')::integer);

-- Admin access policy
CREATE POLICY admin_access ON bookings
  FOR ALL TO admin_user
  USING (true);

-- User profile access policy
CREATE POLICY user_profile_access ON users
  FOR ALL TO authenticated_user
  USING (id = current_setting('app.current_user_id')::integer);

-- Audit log access policy (admin only)
CREATE POLICY audit_log_admin_only ON audit_log
  FOR ALL TO admin_user
  USING (true);

```

### Application-Level Security Integration:

```

// Middleware for setting user context
const setUserContext = (req, res, next) => {
  if (req.user && req.user.id) {
    req.dbClient.query('SET app.current_user_id = $1', [req.user.id]);
  }
  next();
};

// Secure booking creation
const createBooking = async (req, res) => {
  const client = await postgresPool.connect();
  try {
    await client.query('BEGIN');

    // Set user context for row-level security
    await client.query('SET app.current_user_id = $1', [req.user.id]);

    const booking = await client.query(`
      INSERT INTO bookings (user_id, trip_code, travel_date, travelers,
total_amount)
      VALUES ($1, $2, $3, $4, $5)
    `);
  } catch (error) {
    await client.query('ROLLBACK');
    res.status(500).json({ error: error.message });
  }
};

```

```

        RETURNING *
    `, [req.user.id, req.body.tripCode, req.body.travelDate,
req.body.travelers, req.body.totalAmount]);

    await client.query('COMMIT');
    res.json(booking.rows[0]);
} catch (error) {
    await client.query('ROLLBACK');
    throw error;
} finally {
    client.release();
}
};

```

## Audit Logging System

### Comprehensive Audit Trail Implementation:

```

-- Audit trigger function
CREATE OR REPLACE FUNCTION audit_trigger_function()
RETURNS TRIGGER AS $$
BEGIN
    IF TG_OP = 'DELETE' THEN
        INSERT INTO audit_log (table_name, record_id, operation, old_values,
user_id, ip_address, user_agent)
        VALUES (TG_TABLE_NAME, OLD.id, TG_OP, row_to_json(OLD),
current_setting('app.current_user_id')::integer,
                current_setting('app.client_ip')::inet,
current_setting('app.user_agent')::text);
        RETURN OLD;
    ELSIF TG_OP = 'UPDATE' THEN
        INSERT INTO audit_log (table_name, record_id, operation, old_values,
new_values, user_id, ip_address, user_agent)
        VALUES (TG_TABLE_NAME, NEW.id, TG_OP, row_to_json(OLD),
row_to_json(NEW),
                current_setting('app.current_user_id')::integer,
current_setting('app.client_ip')::inet,
                current_setting('app.user_agent')::text);
        RETURN NEW;
    ELSIF TG_OP = 'INSERT' THEN
        INSERT INTO audit_log (table_name, record_id, operation, new_values,
user_id, ip_address, user_agent)
        VALUES (TG_TABLE_NAME, NEW.id, TG_OP, row_to_json(NEW),
current_setting('app.current_user_id')::integer,
                current_setting('app.client_ip')::inet,
current_setting('app.user_agent')::text);

```

```

        RETURN NEW;
    END IF;
    RETURN NULL;
END;
$$ LANGUAGE plpgsql;

-- Apply audit triggers
CREATE TRIGGER bookings_audit_trigger
    AFTER INSERT OR UPDATE OR DELETE ON bookings
    FOR EACH ROW EXECUTE FUNCTION audit_trigger_function();

CREATE TRIGGER users_audit_trigger
    AFTER INSERT OR UPDATE OR DELETE ON users
    FOR EACH ROW EXECUTE FUNCTION audit_trigger_function();

```

## Performance Optimization

### Strategic Indexing Implementation:

```

-- Composite indexes for common query patterns
CREATE INDEX idx_bookings_user_date ON bookings(user_id, travel_date);
CREATE INDEX idx_bookings_status_date ON bookings(status, booking_date);
CREATE INDEX idx_users_email_active ON users(email, is_active);
CREATE INDEX idx_audit_log_table_timestamp ON audit_log(table_name,
timestamp);

-- Partial indexes for active records
CREATE INDEX idx_active_users ON users(id) WHERE is_active = true;
CREATE INDEX idx_active_bookings ON bookings(id) WHERE status !=
'cancelled';

```

### Connection Pooling Configuration:

```

// PostgreSQL connection pool
const postgresPool = new Pool({
    host: process.env.POSTGRES_HOST,
    port: process.env.POSTGRES_PORT,
    database: process.env.POSTGRES_DB,
    user: process.env.POSTGRES_USER,
    password: process.env.POSTGRES_PASSWORD,
    max: 20, // Maximum number of clients in the pool
    idleTimeoutMillis: 30000, // Close idle clients after 30 seconds
    connectionTimeoutMillis: 2000, // Return an error after 2 seconds if
connection could not be established
    ssl: process.env.NODE_ENV === 'production' ? { rejectUnauthorized: false
} : false

```

```

});

// MongoDB connection with optimization
const MongoClient = new MongoClient(process.env.MONGODB_URI, {
  maxPoolSize: 10, // Maintain up to 10 socket connections
  serverSelectionTimeoutMS: 5000, // Keep trying to send operations for 5
seconds
  socketTimeoutMS: 45000, // Close sockets after 45 seconds of inactivity
  bufferMaxEntries: 0, // Disable mongoose buffering
  bufferCommands: false // Disable mongoose buffering
});

```

## Backup and Recovery Systems

### Automated Backup Implementation:

```

// Backup service implementation
class BackupService {
  constructor() {
    this.postgresConfig = {
      host: process.env.POSTGRES_HOST,
      port: process.env.POSTGRES_PORT,
      database: process.env.POSTGRES_DB,
      user: process.env.POSTGRES_USER,
      password: process.env.POSTGRES_PASSWORD
    };
  }

  async createPostgreSQLBackup() {
    const timestamp = new Date().toISOString().replace(/[:.]/g, '-');
    const backupFile = `backup_${timestamp}.sql`;

    const command = `pg_dump -h ${this.postgresConfig.host} -p
${this.postgresConfig.port} -U ${this.postgresConfig.user} -d
${this.postgresConfig.database} -f ${backupFile}`;

    return new Promise((resolve, reject) => {
      exec(command, (error, stdout, stderr) => {
        if (error) {
          reject(error);
        } else {
          resolve(backupFile);
        }
      });
    });
  }
}

```

```

async createMongoDBBackup() {
  const timestamp = new Date().toISOString().replace(/[:.]/g, '-');
  const backupDir = `mongodb_backup_${timestamp}`;

  const command = `mongodump --uri="${process.env.MONGODB_URI}"
--out=${backupDir}`;

  return new Promise((resolve, reject) => {
    exec(command, (error, stdout, stderr) => {
      if (error) {
        reject(error);
      } else {
        resolve(backupDir);
      }
    });
  });
}

async scheduleBackups() {
  // Daily full backup
  cron.schedule('0 2 * * *', async () => {
    try {
      await this.createPostgreSQLBackup();
      await this.createMongoDBBackup();
      console.log('Daily backup completed successfully');
    } catch (error) {
      console.error('Backup failed:', error);
    }
  });
}
}

```

## Challenges Faced and Solutions

### Challenge 1: Data Consistency Across Databases

- **Problem:** Ensuring data consistency between PostgreSQL and MongoDB
- **Solution:** Implemented event-driven synchronization with transaction logs
- **Learning:** Understanding eventual consistency and distributed system challenges

### Challenge 2: Performance Impact of Security Policies

- **Problem:** Row-level security policies could impact query performance
- **Solution:** Optimized indexes and query patterns to work with security policies
- **Learning:** Balancing security requirements with performance considerations



### Challenge 3: Backup and Recovery Complexity

- **Problem:** Coordinating backups across different database systems
- **Solution:** Implemented coordinated backup strategies with point-in-time recovery
- **Learning:** Understanding disaster recovery planning and implementation

## Course Outcomes Demonstrated

### Outcome 3: Computing Solutions

This enhancement demonstrates my ability to design and evaluate computing solutions using appropriate database principles. The polyglot persistence architecture showcases my understanding of different database paradigms and their trade-offs.

#### Specific Evidence:

- Strategic data distribution based on data characteristics
- ACID compliance for transactional data
- Flexible schema for content management
- Performance optimization through indexing

### Outcome 5: Security Mindset

The comprehensive security implementation demonstrates my ability to anticipate and mitigate database security vulnerabilities. The row-level security and audit logging provide defense-in-depth protection.

#### Specific Evidence:

- Row-level security preventing unauthorized data access
- Comprehensive audit logging for compliance
- Data encryption at rest and in transit
- SQL injection prevention through prepared statements

### Outcome 4: Technical Implementation

The use of modern database technologies and enterprise practices demonstrates my ability to implement industry-standard solutions. The integration of multiple database systems showcases technical proficiency.

#### Specific Evidence:

- PostgreSQL and MongoDB integration
- Connection pooling for scalability
- Automated backup and recovery systems
- Performance monitoring and optimization

## Reflection on Learning and Growth

This enhancement process was particularly valuable in developing my understanding of enterprise database systems. The experience of designing and implementing a polyglot persistence architecture taught me the importance of choosing the right tool for each specific use case.

The security implementation work required deep understanding of database security principles and their practical application. Learning to implement row-level security and audit logging has prepared me for working with sensitive data in professional environments.

The performance optimization work taught me the importance of understanding query patterns and access patterns when designing database schemas. The experience of implementing strategic indexing and connection pooling has given me practical skills for optimizing database performance.

Working with backup and recovery systems introduced me to disaster recovery planning and business continuity considerations. This knowledge is crucial for maintaining data availability in production environments.

## Future Improvements and Considerations

Several areas for future enhancement include:

1. **Read Replicas:** Implementing read replicas for improved read performance
2. **Sharding:** Implementing horizontal partitioning for extreme scalability
3. **Data Lake Integration:** Adding data warehouse capabilities for analytics
4. **Real-time Synchronization:** Implementing change data capture for real-time updates

This enhancement demonstrates my ability to design, implement, and secure complex database systems that meet enterprise requirements. The skills developed through this process are directly applicable to professional database administration and software development roles requiring advanced data management capabilities.